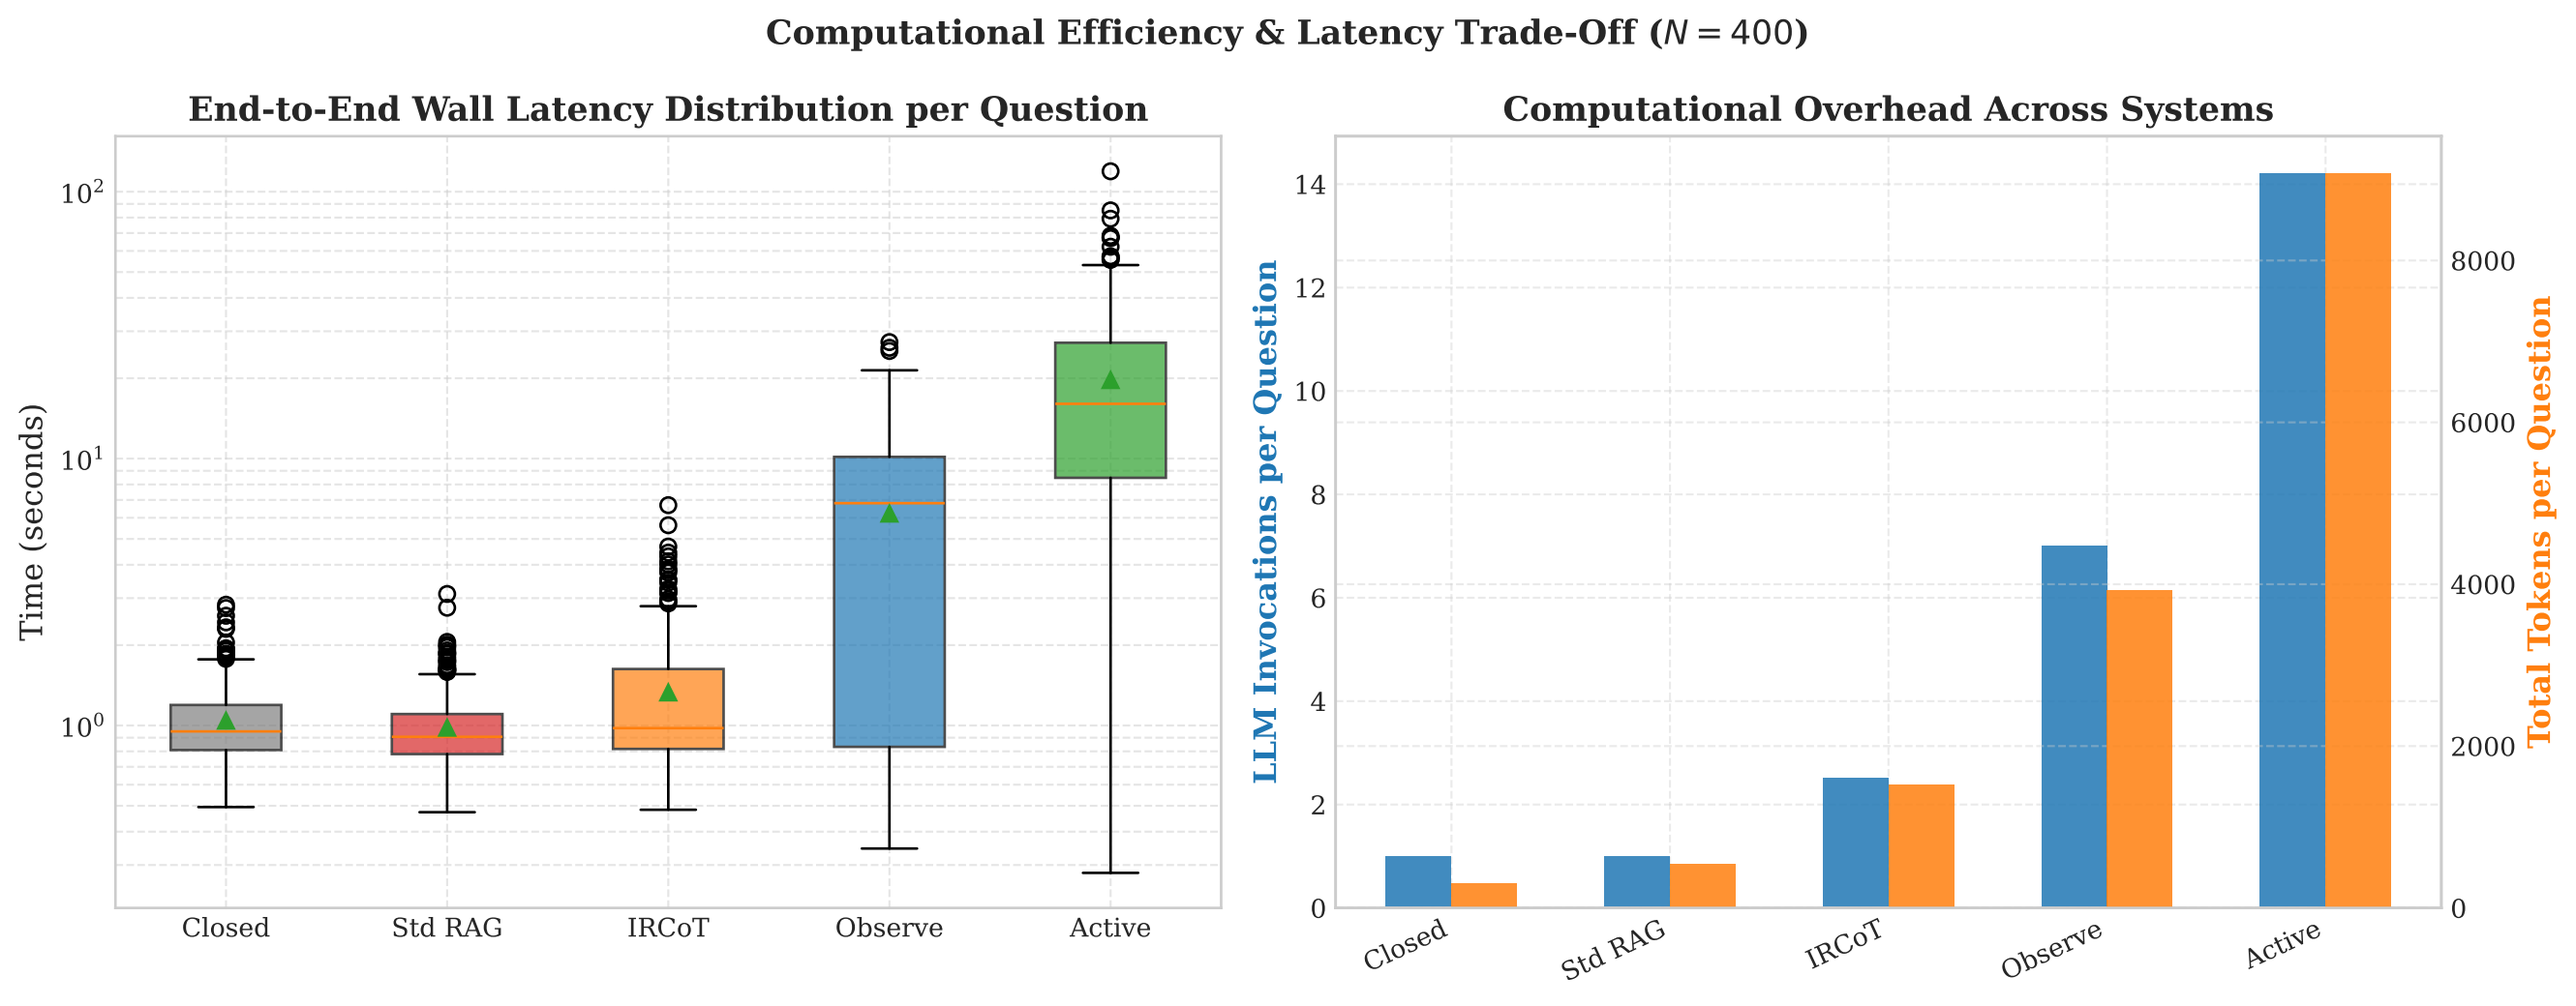
Computational Efficiency & Latency Trade-Off ($N = 400$)
End-to-End Wall Latency Distribution per Question
Time (seconds)
10^2
10^1
10^0
Closed
Std RAG
IRCoT
Observe
Active
Computational Overhead Across Systems
LLM Invocations per Question
Total Tokens per Question
0
2
4
6
8
10
12
14
0
2000
4000
6000
8000
Closed
Std RAG
IRCoT
Observe
Active