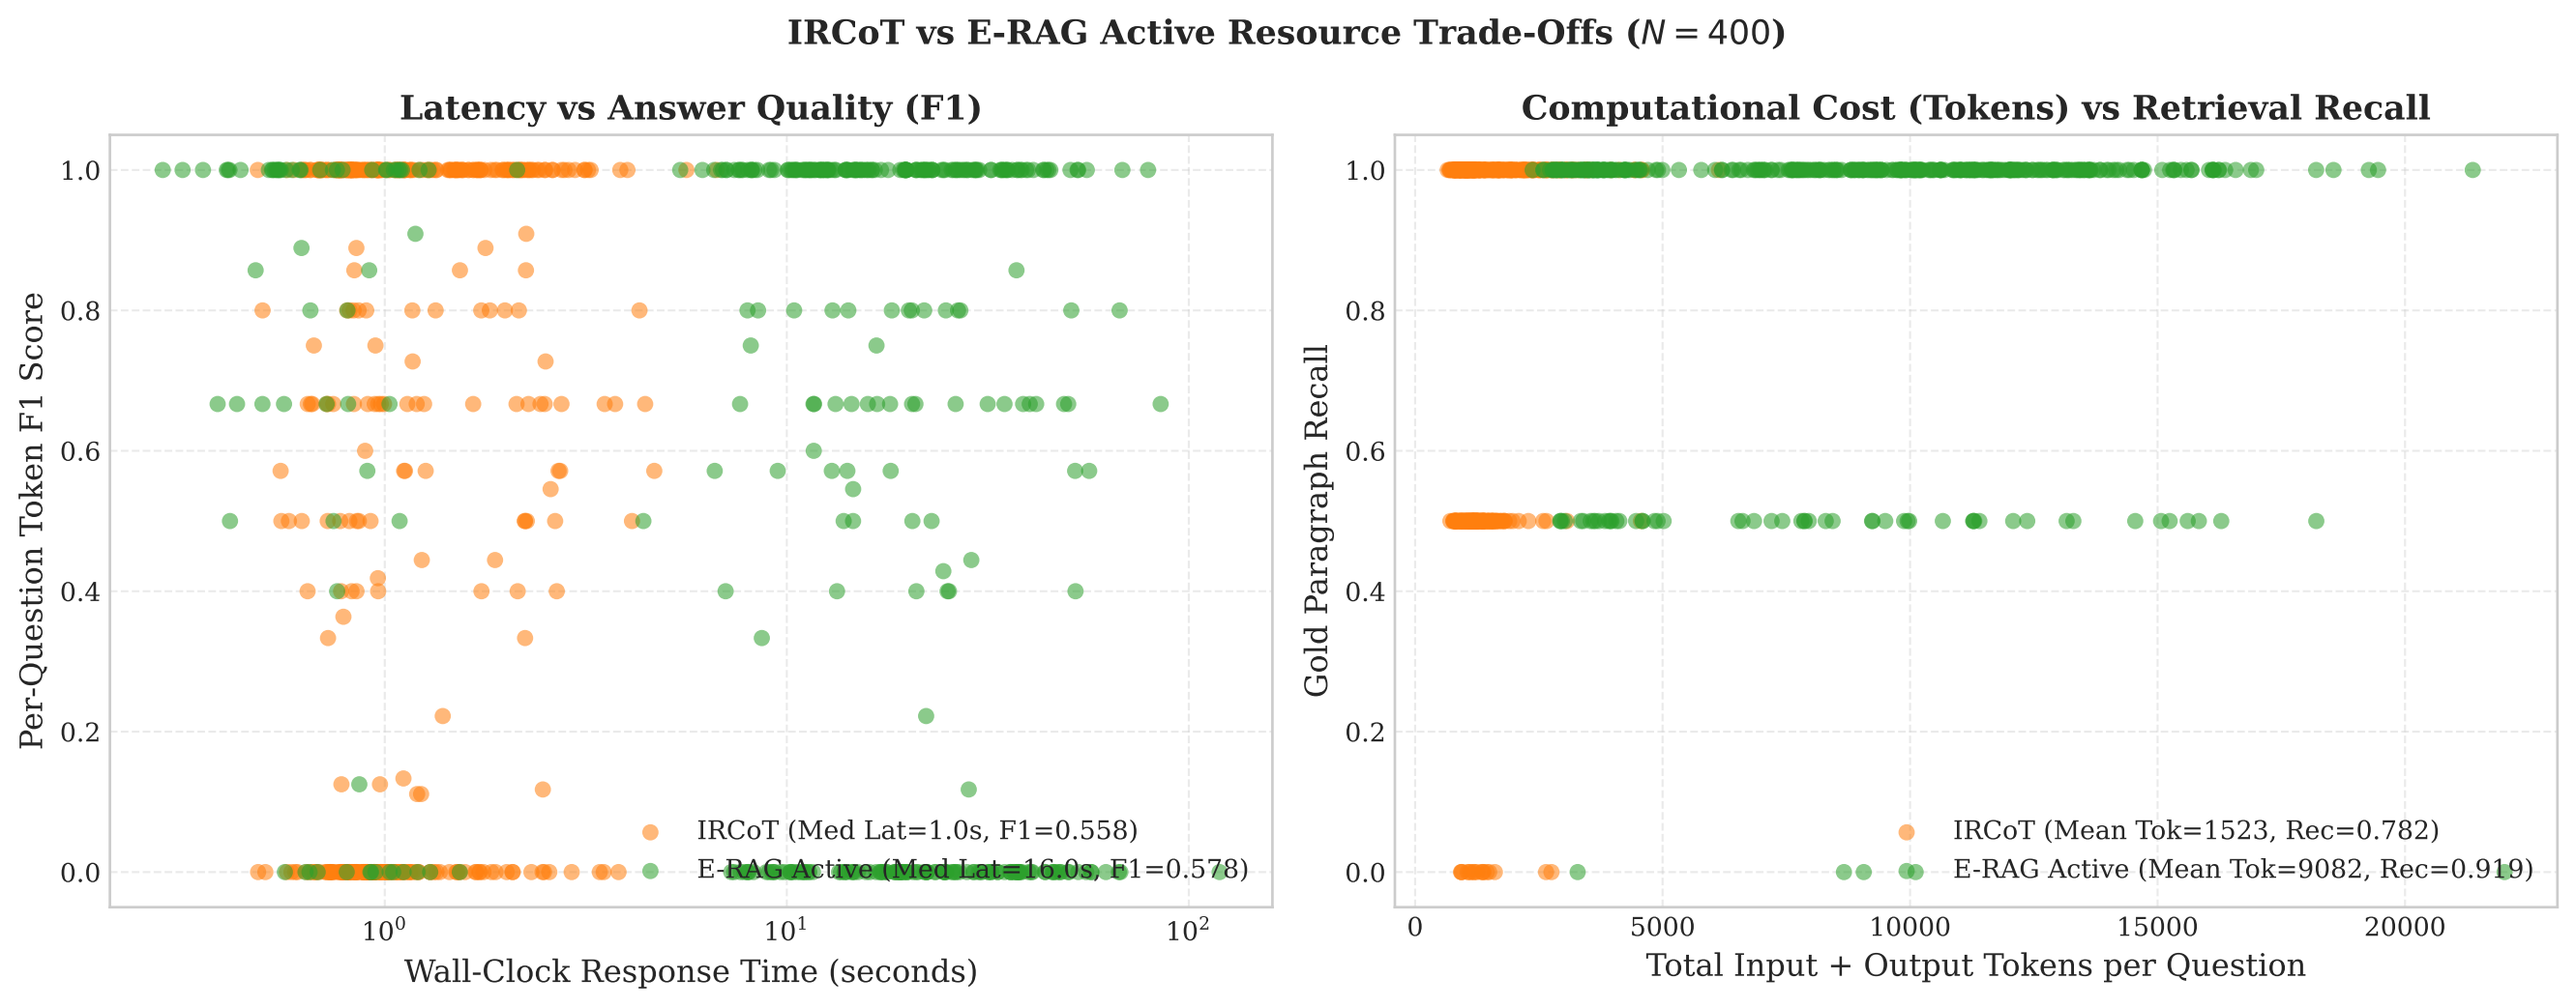
IRCoT vs E-RAG Active Resource Trade-Offs ($N = 400$)
Latency vs Answer Quality (F1)
Per-Question Token F1 Score
Wall-Clock Response Time (seconds)
IRCoT (Med Lat=1.0s, F1=0.558)
E-RAG Active (Med Lat=16.0s, F1=0.578)
Computational Cost (Tokens) vs Retrieval Recall
Gold Paragraph Recall
Total Input + Output Tokens per Question
IRCoT (Mean Tok=1523, Rec=0.782)
E-RAG Active (Mean Tok=9082, Rec=0.919)